arXiv:2306.07902v1 [cs.CL] 13 Jun 2023

Massively Multilingual Corpus of Sentiment Datasets and Multi-faceted Sentiment Classification Benchmark

Łukasz Augustyniak
WUST (Wrocław University of Science and Technology)
lukasz.augustyniak@pwr.wroc.pl

Szymon Woźniak
Brand24 AI

Marcin Gruza
Brand24 AI, WUST

Piotr Gramacki
Brand24 AI, WUST

Krzysztof Rajda
Brand24 AI, WUST

Mikołaj Morzy
Poznań University of Technology

Tomasz Kajdanowicz
WUST

Abstract

Despite impressive advancements in multilingual corpora collection and model training, developing large-scale deployments of multilingual models still presents a significant challenge. This is particularly true for language tasks that are culture-dependent. One such example is the area of multilingual sentiment analysis, where affective markers can be subtle and deeply ensconced in culture. This work presents the most extensive open massively multilingual corpus of datasets for training sentiment models. The corpus consists of 79 manually selected datasets from over 350 datasets reported in the scientific literature based on strict quality criteria. The corpus covers 27 languages representing 6 language families. Datasets can be queried using several linguistic and functional features. In addition, we present a multi-faceted sentiment classification benchmark summarizing hundreds of experiments conducted on different base models, training objectives, dataset collections, and fine-tuning strategies.

1 Introduction

Consider a hotel booking service that allows its customers to post reviews. You have found just the perfect accommodation to stay for a couple of days with your family, but you browse through the reviews section of the website to check the experiences of former guests. Suddenly, you encounter a review in Polish: "*hotel jak hotel, mogło być gorzej*." This review has the following sentiment scores[1]: $s_{neg} = 0.44, s_{neu} = 0.44, s_{pos} = 0.12$. Intrigued by the ambiguity of scores, you translate the review into English: "*hotel like a hotel, all in all, it could have been worse*," which is scored as $s_{neg} = 0.80, s_{neu} = 0.16, s_{pos} = 0.04$. Apparently, the stereotypically pessimistic Polish outlook on life gets lost in translation. The next review is written in Czech: "*ok, ale nic zajímavého*" with scores $s_{neg} = 0.32, s_{neu} = 0.54, s_{pos} = 0.14$. After translating into English ("*ok, but nothing interesting*") the sentiment classification model scores the review as negative ($s_{neg} = 0.50, s_{neu} = 0.37, s_{pos} = 0.13$). After your stay, you decide to add a very positive review of the hotel ("*it was a killer place to stay*", $s_{neg} = 0.03, s_{neu} = 0.05, s_{pos} = 0.92$). You would be very surprised to learn that the Czechs would be quite confused about your opinion ("*to bylo vražedné*

[1]Sentiment scores in this paragraph are produced by the multilingual `cardiffnlp/twitter-xlm-roberta-base-sentiment` model [9]

Preprint. Under review.

místo k pobytu", $s_{neg} = 0.51, s_{neu} = 0.09, s_{pos} = 0.39$), while the Poles would stay away from the hotel at all costs ("*to było zabójcze miejsce na pobyt*", $s_{neg} = 0.78, s_{neu} = 0.09, s_{pos} = 0.13$).

multilingual services become ubiquitous in the modern global economy. As more websites begin to offer automatic translation of content, users do not bother to express themselves in the *lingua franca* of the Web, writing instead in their native languages. Despite impressive advancements in automatic translation, many NLP tasks are still difficult in the multilingual setting. And sentiment classification is among the most challenging. The expression of sentiment is highly culture-dependent [28]. The emotional valence of individual words, the presence of specific phrasemes, and the expectations around sentiment values make sentiment classification across languages a demanding task.

Models performing sentiment classifications have to cope with two independent sources of variance in the input data: cultural expressions of sentiment and errors in automatic translations. In addition, the productization of sentiment classification leads to several engineering choices which influence the efficiency of the model:

- *single multilingual model vs. dedicated monolingual models*: deploying a single model in a production environment is much easier than orchestrating an ensemble of models,
- *training vs. fine-tuning*: sentiment classification model can be trained from scratch, both in the multi- and monolingual regimes, and the alternative is the fine-tuning of a general-purpose pre-trained language model with sentiment data,
- *transfer learning between domains*: a sentiment classification model trained for a specific domain (e.g., book reviews) can be transferred to another domain (e.g., hotel reviews) under the assumption that sentiment expressions in a given language remain independent of the subject of sentiment,
- *transfer learning between languages*: one may hypothesize that related languages can utilize similar sentiment expression mechanisms regarding grammar, punctuation, and vocabulary, so it is possible to use training data in a language other than the target language to train or fine-tune sentiment classification models.

Working with multilingual datasets and models opens another opportunity. Most NLP research focuses on 20 major languages. Many languages native to significant human populations are not adequately studied. Although the term *low-resource language* is not precisely defined and can be understood in terms of computerization, privilege, the abundance of resources, or density [43], the existence of a large chasm between languages with respect to linguistic resources is apparent. Our work aims at supporting low-resource languages in performant sentiment classification. Finally, there is a lively debate in the scientific community about the inherent ability of neural models to handle general linguistic phenomena. This paper tries to answer this question in the domain of sentiment classification.

Contribution presented in this paper is threefold:

- *multilingual corpus*: we present the largest multilingual collection of sentiment classification datasets consisting of 79 high-quality datasets covering 27 languages. The collection of the corpus and detailed descriptions of collected datasets are presented in Section 3.
- *multi-faceted benchmark*: the dataset is supplemented with a benchmark containing detailed run statistics of hundreds of experiments representing different training and testing scenarios. All details relevant to the benchmark are presented in Section 4.
- *library for dataset access*: all datasets in the multilingual corpus are publicly available via a library compatible with the HuggingFace library, along with the ability to filter and select datasets, verify their licenses, etc[2].

2 Linguistic typology

The similarity of languages and the main aspects of their differences is the field of study of language typology. The differences between languages can be phonological (differences in sounds used by languages), syntactic (differences in language structures), lexical (differences in vocabulary),

[2] `https://huggingface.co/datasets/Brand24/mms`

and theoretical (differences characterized as general properties of languages). Linguistic typology analyzes the current state of languages and is often contrasted with genealogical linguistics. The latter is concerned with historical relationships between languages established via historical records or with the help of comparative linguistics. The main focus of genealogical linguistics resolves around *language families* and *language genera*. The term *language family* refers to a group of languages sharing pedigree from a common ancestral language (the *proto-language*). As of today, linguists define over 7000 languages categorized into 150 families. The largest families of languages include Indo-European, Sino-Tibetan, Turkic, Afro-Asiatic, Nilo-Saharan, Niger-Congo, and Eskimo-Aleut. Main families are further divided into branches called *genera*. Examples of genera within the Indo-European family of languages include Slavic, Romance, Germanic, and Indic. This division of language families into genera closely follows the genetic family of humankind as attested by DNA similarity. Some language families do not produce distinctive genera but form dialect continua defined by mutual intelligibility. Finally, an important concept is that of a *sprachbund*, a geographical area occupied by languages sharing linguistic features. The linguistic similarities within a sprachbund result from cultural exchange and geographical contact rather than by chance or common origin.

Languages can be described using hundreds of linguistic features. World Atlas of Language Structures [24] lists almost 200 different features. Since our work focuses on sentiment classification, we select 10 features that seem to be the most relevant to the task of sentiment expression. These features are:

1. *definite article*: the morpheme associated with nouns and used to code the uniqueness or definiteness of a concept; almost half of the languages do not use the definite article.
2. *indefinite article*: the morpheme used together with nouns to signal that the related concept is unknown to the hearer; half of the languages do not use the indefinite article, some languages use a separate article, and some use the numeral "one" as the indefinite article.
3. *number of cases*: morphological cases are a common way to express various relationships between words; human languages vary greatly in the number of cases used by a language.
4. *order of subject, verb, and object*: some languages have the strict ordering of words, while languages that convey semantics through inflection may be much looser with the ordering; half of all languages use the SOV (subject-object-verb) ordering, one third uses the SVO (subject-verb-object) ordering, and a small fraction of languages use the remaining VSO, VOS, OVS, and OSV orderings. Interestingly, around 13% of world languages do not have any fixed word order.
5. *negative morphemes*: negative morpheme is used to signal clausal negation in declarative sentences; this is usually achieved using a negative affix or a negative particle.
6. *polar questions*: a polar question is a question with only yes/no answers; these questions can be built using question particles, interrogative morphology, or intonation only.
7. *position of the negative morpheme*: languages differ by the position of the negative morpheme in relation to subjects and objects, with many variants such as SNegVO, NegSVO, SVNegO, obligatory and optional double negations, etc.
8. *prefixing vs. suffixing*: languages differ significantly in their use of prefixes versus suffixes in inflectional morphology.
9. *coding of nominal plurals*: two major types of plural coding are present in languages, either by changing the morphological form of the noun or by using a plurality indicator morpheme somewhere in the noun phrase.
10. *grammatical genders*: there is significant variability among languages with respect to the number of grammatical genders, some languages do not use the concept at all, and some languages may have 5 or more grammatical genders.

All language features mentioned above are available as filtering features in our library. Thus, when training a sentiment classifier using our dataset, one may download different facets of the collection. For instance, one can download all datasets in Slavic languages in which polar questions are formed using the interrogative word order (Listing 1) or download all datasets from the Afro-Asiatic language family with no morphological case-making (Listing 2).

3 Datasets

3.1 Quality criteria

The initial pool of sentiment datasets has been gathered using extensive search and consisted of 345 datasets found through Google Scholar, GitHub repositories, and the HuggingFace datasets library. This initial pool of datasets has been manually filtered based on the following set of quality assurance criteria:

1. *strong annotations*: we have rejected datasets containing weak annotations (e.g., datasets with labels based on emoji occurrence or generated automatically through classification by machine learning models) due to an extensive amount of noise [51].
2. *well-defined annotation protocol*: we have rejected datasets without sufficient information about the annotation protocol (e.g., whether annotation was manual or automatic, number of annotators) to avoid merging datasets with contradicting annotation instructions.
3. *numerical ratings*: we have accepted datasets with numerical ratings, mapping Likert-type 5-point scales into three class sentiment labels as follows: ratings 1 and 2 were mapped to *negative*, rating 3 was mapped to *neutral*, and ratings 4 and 5 were mapped to *positive*.
4. *three classes only*: we have rejected datasets annotated with binary sentiment labels as their performance in three class settings was unsatisfactory.
5. *monolingual datasets*: when a dataset contained samples in multiple languages, we opted to divide it into independent datasets in constituent languages.

3.2 Pre-processing of datasets

Despite quality assurance criteria described in Section 3.1, the datasets still contained conflicting entries, i.e., duplicated records with different sentiment labels. We have cross-referenced all datasets to identify conflicts and have made the data coherent using majority-label voting. Finally, labeling and rating schemes of all datasets have been mapped to a 3-class scheme with *negative*, *neutral*, and *positive* labels only. For datasets with emotional annotations, we mapped positive emotions (joy, happiness) into positive sentiment and negative emotions (fear, sadness, disgust, anger) into negative sentiment. Texts with ambiguous emotions like anticipation and surprise were discarded. This pre-processing pipeline resulted in 79 datasets containing $6\,164\,762$ text samples. Most of the datasets are in English ($2\,330\,486$ samples across 17 datasets), Arabic ($932\,075$ samples across 9 datasets), and Spanish ($418\,712$ samples across 5 datasets). The datasets represent four different domains: social media (44 datasets), reviews (24 datasets), news (5 datasets), and others (6 datasets). In addition, all datasets were processed by the `cleanlab` library to produce a self-confidence label-quality score for each data point.

In order to normalize the testing of hundreds of configurations of the benchmark, we have prepared a separate internal dataset with manual annotations. This dataset covers multiple domains and consists of $3\,911$ samples in Polish and English annotated independently by 3 annotators with the majority label selection. The inter-rater agreement was $\kappa = 0.665$ (Cohen's kappa) and $\alpha = 0.666$ (Krippendorff's alpha). We use the internal dataset as the universal validation dataset in the experiments presented in the benchmark.

4 Multi-faceted benchmark

As we have explained in Section 1, the deployment of a multilingual sentiment classifier can be evaluated using several criteria leading to different architecture choices. Thus, we do not publish a single benchmark, but we aggregate the results along several dimensions. The benchmark available at `https://huggingface.co/spaces/Brand24/mms_benchmark`, allows to compare models according to:

- *number of languages*: multilingual models with monolingual models,
- *training procedure*: models trained from scratch vs. fine-tuning,
- *domain language*: language to which a model is applied,

Table 1: Summary of the corpus. Categories: N - news, R - reviews, SM - social media, O - other

	#datasets	category				#samples			mean	
		N	R	SM	O	NEG	NEU	POS	#words	#chars
English	17	3	4	6	4	304,939	290,823	1,734,724	62	339
Arabic	9	0	4	4	1	138,899	192,774	600,402	52	289
Spanish	5	0	3	2	0	108,733	122,493	187,486	26	150
Chinese	2	0	2	0	0	117,967	69,016	144,719	60	80
German	6	0	1	5	0	104,667	100,071	111,149	26	171
Polish	4	0	2	2	0	77,422	62,074	97,192	19	123
French	3	0	1	2	0	84,187	43,245	83,199	28	159
Japanese	1	0	1	0	0	83,982	41,979	83,819	61	101
Czech	4	0	2	2	0	39,674	59,200	97,413	34	212
Portuguese	4	0	0	4	0	56,827	55,165	45,842	11	63
Slovenian	2	1	0	1	0	33,694	50,553	29,296	41	269
Russian	2	0	0	2	0	31,770	48,106	31,054	11	70
Croatian	2	1	0	1	0	19,757	19,470	38,367	17	116
Serbian	3	0	2	1	0	25,089	32,283	18,996	44	269
Thai	2	0	1	1	0	9,326	28,616	34,377	22	381
Bulgarian	1	0	0	1	0	13,930	28,657	19,563	12	86
Hungarian	1	0	0	1	0	8,974	17,621	30,087	11	83
Slovak	1	0	0	1	0	14,431	12,842	29,350	13	98
Albanian	1	0	0	1	0	6,889	14,757	22,638	13	91
Swedish	1	0	0	1	0	16,266	13,342	11,738	14	94
Bosnian	1	0	0	1	0	11,974	11,145	13,064	12	76
Urdu	1	0	0	0	1	5,239	8,585	5,836	13	69
Hindi	1	0	0	1	0	4,992	6,392	5,615	26	128
Persian	1	0	1	0	0	1,602	5,091	6,832	21	104
Italian	2	0	0	2	0	4,043	4,193	3,829	16	103
Hebrew	1	0	0	1	0	2,279	243	6,097	22	110
Latvian	1	0	0	1	0	1,378	2,618	1,794	20	138

- *data modality*: news, reviews, social media,
- *knowledge transfer*: transfer between languages, transfer between domains.

Table 2 presents the list of models included in the benchmark. For each model, we include the number of parameters and languages used in pre-training and the base model. The results presented in the benchmark reflect three possible scenarios of model deployment. A pre-trained model can be used to generate text representation only. In this scenario, denoted *head-linear* (HL), a model serves as a feature extractor followed by a small linear classification head. In the second scenario, the linear classification head is replaced by a BiLSTM classifier operating on features extracted by the pre-trained model. We refer to this scenario as *head-bilstm* (HB). Finally, each pre-trained transformer-based model (with the exception of mUSE-transformer) has been fine-tuned to the sentiment classification task. We refer to this scenario as *fine-tuning* (FT).

The hyperparameters of models included in the benchmark are as follows. The hidden size of the model was set to the embedding size of each model when used in HL and FT scenarios. The hidden size of the HB scenario was set to $h = 32$. The learning rate for the HL scenario was $\eta = 1 \times 10^{-3}$, fine-tuning used $\eta = 1 \times 10^{-5}$, and HB scenario used $\eta = 5 \times 10^{-3}$. The batch size in HL and HB scenarios was $b = 200$ and $b = 6$ for fine-tuning. Dropout for HB was $d = 0.5$ and $d = 0.2$ for other scenarios. Training took 5 epochs for fine-tuning and 2 epochs for HL and HB scenarios (beyond 2 epochs most models started to overfit). The models are evaluated using the traditional F_1 score computed on three levels: the entire dataset, averaged over all datasets, and the internal dataset.

We performed our experiments using Python 3.9 and PyTorch (1.8.1) (and Tensorflow (2.3.0) for original mUSE). Our experimental setup consists of Intel(R) Xeon(R) CPU E5-2630 v4 @ 2.20GHz and Nvidia Tesla V100 16GB.

Table 2: Models included in the benchmark

Model	#params	#langs	base	reference
mT5	277M	101	T5	Xue et al. [82]
LASER	52M	93	BiLSTM	Artetxe and Schwenk [5]
mBERT	177M	104	BERT	Devlin et al. [23]
MPNet	278M	53	XLM-R	Reimers and Gurevych [56]
XLM-R-dist	278M	53	XLM-R	Reimers and Gurevych [56]
XLM-R	278M	100	XLM-R	Conneau et al. [20]
LaBSE	470M	109	BERT	Feng et al. [27]
DistilmBERT	134M	104	BERT	Sanh et al. [64]
mUSE-dist	134M	53	DistilmBERT	Reimers and Gurevych [56]
mUSE-transformer	85M	16	transformer	Yang et al. [84]
mUSE-cnn	68M	16	CNN	Yang et al. [84]

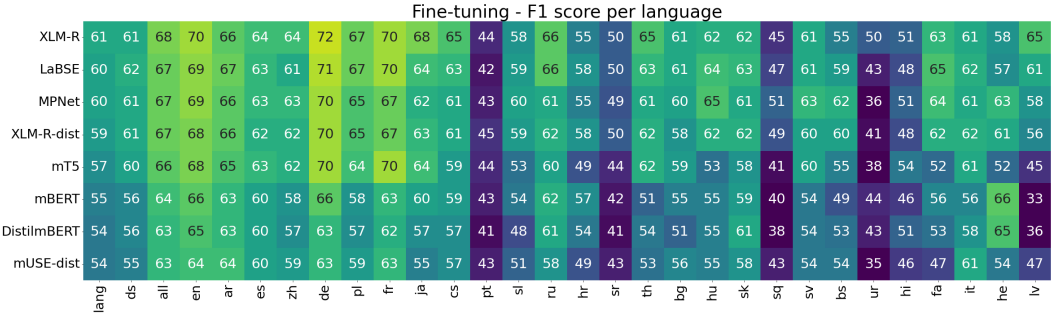

Figure 1: Detailed results of models' comparison.

Figure 2: Legend: **lang** - averaged by all languages, **ds** - averaged by dataset, **ar** - Arabic, **bg** - Bulgarian, **bs** - Bosnian, **cs** - Czech, **de** - German, **en** - English, **es** - Spanish, **fa** - Persian, **fr** - French, **he** - Hebrew, **hi** - Hindi, **hr** - Croatian, **hu** - Hungarian, **it** - Italian, **ja** - Japanese, **lv** - Latvian, **pl** - Polish, **pt** - Portuguese, **ru** - Russian, **sk** - Slovak, **sl** - Slovenian, **sq** - Albanian, **sr** - Serbian, **sv** - Swedish, **th** - Thai, **ur** - Urdu, **zh** - Chinese.

Here we present one result available through the benchmarks. Due to space constraints, we move the presentation of other case studies to Appendix B. Figure 2 contains F_1 scores of all models fine-tuned on all datasets available for a given language. Interestingly, we see little variance in performance between models for high-resource languages and significant deterioration of performance for low-resource languages. An outlier is the aforementioned Portuguese, where the performance is caused by the lack of data points representing news and reviews. Figure 2 directly shows the usefulness of our benchmark. When designing sentiment classification solutions in Spanish, the choice of model is secondary (models' performances are very similar), and other model features can be considered (such as ease of deployment, cost of inference, and memory requirements). However, if sentiment classification is to be applied to Latvian, the performance difference between models can be as high as 32 pp depending on the choice of the model.

5 Related work

5.1 Multilingual text representations

One of the foundational discoveries in multilingual presentation learning was the fact that latent vector spaces seemed to encode very similar word relationships across a wide spectrum of languages. Needless to say, first monolingual static word embeddings were quickly followed by multilingual word embeddings which provided the basis for multilingual text representations [59]. A popular solution was to use pre-trained monolingual embeddings, such as `word2vec` [45] and align them via linear transformations using parallel multilingual dictionaries [46]. Another approach advocated for joint learning of multilingual word embeddings on pseudo-bilingual datasets, where tokens in one language would be randomly translated to another language [81].

Static multilingual embeddings were superseded by contextual embeddings produced by language models such as BiLSTM [5] and various Transformers [27, 20, 23, 82, 84]. In the case of these more complex architectures, the multilingual capabilities of language models resulted from specific objective functions used during training. These objectives "pushed" the models toward universal multilingual representations by forcing the models to perform machine translation [5], translation language modeling (TLM) [20, 18], or translation ranking [27, 83].

The evaluation of the quality of multilingual text representation is not trivial. Cross-lingual and multilingual tasks are actively developed to foster the development of multilingual models. Examples of such cross-lingual and multilingual tasks include cross-lingual natural language inference [19], question answering [39], named entity recognition [76, 77] or parallel text extraction [86, 85]. For example, results of comparison of LASER, mBERT, and XLM in tasks of named entity recognition and part-of-speech tagging in zero-shot settings suggests that LASER outperforms the latter methods in the case of knowledge transfer [78]. Another important benchmark is XTREME [31], designed for testing the abilities of cross-lingual transfer across 40 languages and 9 tasks. Despite its massive character, XTREME lacks benchmarking task of sentiment analysis. Also, only mBERT, XLM, XLM-R, and MMTE are used as baseline models.

By far, the most popular pre-trained multilingual language model is Multilingual BERT (mBERT) [23]. It has been used in several cross-lingual studies, for instance, in zero-shot knowledge transfer between Slovenian and Croatian languages [54], exploring code-switching on Spanglish and Hinglish [53], and building hierarchical architecture for zero-shot setting [65]. The properties and limitations of mBERT have been extensively studied. One of the studies showed that mBERT is not learning a joint representation of languages. Rather, it partitions the representation space between languages [69]. The authors used the Projection Weighted Canonical Correlation Analysis (PWCCA) to analyze how translations of the same sentence are represented in mBERT layers. The correlations were stable in pre-trained and fine-tuned models, with the effect being more pronounced in deeper layers of the model. The hierarchical structure induced by correlations was similar to the structure produced by genealogical linguistics.

Another interesting finding was that mBERT encodes language-specific information within the parameter space, and this language component is not removed by fine-tuning [40]. The language-specific component of mBERT can be removed by estimating the centroid of the language (the mean of mBERT embeddings of a given language vocabulary) and subtracting this centroid from representations produced in the language. The existence of the language component has been proven in multiple tasks: language classification, language similarity, sentence retrieval, word alignment, and machine translation.

Several works tested the cross-lingual abilities of mBERT as compared to monolingual models, finding its performance on low-resource languages inferior to monolingual models [80]. One experiment used a bilingual version of mBERT and trained it in multiple configurations to test the influence of the linguistic relationship between the source and the target language, the network architecture, and the input and learning objective [34]. The authors found that structural similarity and depth of a model are the most significant factors behind mBERT's cross-lingual performance. On the other hand, multi-head attention was not particularly important. Regarding the low-resource languages, [80] compared mBERT with baseline models on tasks of named entity recognition, universal part-of-speech tagging, and universal dependency parsing. The authors focused on the comparison between low- and high-resource languages (determined by the size of the Wikipedia dump in each language). mBERT was found to work well on high-resource languages but under-performed on low-resource languages. Finally, Liu et al. [42] analyzed the relationship between the contextual aspect of mBERT and the training dataset size in the context of multilingual datasets. The authors compared contextual mBERT embeddings with non-contextual models of Word2Vec and GloVe. As it turns out, mBERT outperforms non-contextual models only when large datasets are available for training or fine-tuning. In the low data regime, as well as when the context window is short (limited input), multilingual mBERT is surpassed by static token embeddings.

5.2 Multilingual sentiment classification

Several surveys present an overview of traditional sentiment classification methods (lexicon-based approaches and shallow models with lexical features engineering) [22, 61]. Recently, deep-learning approaches became more prevalent. Attia et al. [6] use convolutional neural networks on word-level

embeddings of texts in English, German and Arabic. The approach is computationally expensive as it requires separate embedding dictionaries for each language. An alternative approach is to use character-level embeddings. Wehrmann et al. [79] trained such a model for binary sentiment classification of English, German, Portuguese, and Spanish tweets. A similar model was used in [4], but the authors applied multi-stage pre-processing, including dictionary checks, soundex algorithm, and lemmatization. Other examples of deep neural models for multilingual sentiment classification include recurrent neural networks supplied with machine translation. Can et al. [15] trained a model on English reviews and evaluated it on machine-translated reviews in Russian, Spanish, Turkish, and Dutch using the Google Translation API and pre-trained GloVe embeddings for English. A similar approach is presented in [35], where the authors used LASER sentence embeddings to train a sentiment classifier on Polish reviews and used this classifier to score reviews translated into other languages.

6 Discussion

The main focus of the paper is the introduction of a massively multilingual large collection of sentiment datasets and the extensive benchmark of model training and validation scenarios. During the construction of the benchmark, we have gathered valuable experiences that we want to share with the community. Our most important observation is that a single multilingual sentiment classifier can perform approximately equally well for all languages. A small group of pre-trained models (XML-R, LaBSE, MPNet) under fine-tuning scenario produces the best results relative to other models. Obviously, the absolute F_1 score across languages differs significantly (due to data scarcity, data quality, or difficulty of samples), but these selected models always produce top classifiers. Secondly, all models perform better under the fine-tuning scenario, but the performance gain varies from model to model. When evaluated on the test dataset, models gained between 4 pp. (mUSE-dist) and 9 pp. (mBERT, DistilmBERT). When tested on the internal datasets, F_1 gains varied between 0 pp (mUSE-dist) and 20 pp. (DistilmBERT), with mT5 and XML-R improving by 17 pp. and 15 pp., respectively. In general, the largest gains can be observed for models trained with the masked language modeling technique (XML-R, mBERT). We also observe that models which produce lower-quality sentence embeddings gain more from fine-tuning.

We also find that bigger models tend to achieve better performance in all data modalities and training scenarios. Of course, there are counterexamples, e.g., mUSE-dist is smaller than mBERT but achieves better performance in the HL scenario for all datasets. This indicates that the size of the model is an important factor for determining its performance, but other factors, like the domain and the type of pretraining task, may also affect the results. Moreover, we observe that the correlation between model size and model performance is weaker after fine-tuning. This means that one may often find a competitive model with similar performance to the state-of-the-art model but significantly smaller and faster for the production environment.

Our final remark is a warning: in our experiments, we have encountered a significant variability in performance conditioned on a particular data split. In one of the cross-validation fine-tuning experiments, we observed a data fold that produced consistently worse results (up to 4 pp.) for all models. This fold was produced by the same random seed and (most probably) consisted of samples with increased difficulty. Since our experiments involved multiple models, multiple training scenarios, multiple datasets, and multiple languages, we had to rely on data subsampling. Conducting experiments on a fixed seed for sample selection could lead to a similar data fold and, consequently, to biased experimental results.

7 Limitations

Despite the fact that our collection is the largest public collection of multilingual sentiment datasets, it still covers only 27 languages. The collection of datasets is highly biased towards the Indo-European family of languages, English in particular. We attribute this bias to the general culture of scientific publishing and its enforcement of English as the primary carrier of scientific discovery. Our work's main potential negative social impact is that the models developed and trained using the provided datasets may still exhibit better performance for the major languages. This could further perpetuate the existing language disparities and inequality in sentiment analysis capabilities across different languages. Addressing this limitation and working towards more equitable representation

and performance across languages is crucial to avoid reinforcing language biases and the potential marginalization of underrepresented languages. The ethical implications of such disparities should be thoroughly discussed and considered.

An important limitation of our dataset collection is a significant variance in sample quality across all datasets and all languages. Figure 3 presents the distribution of self-confidence label-quality score for each data point computed by the `cleanlab` [51]. The distribution of quality is skewed in favor of popular languages, with low-resource languages suffering from data quality issues. A related limitation is caused by an unequal distribution of data modalities across languages. For instance, our benchmark clearly shows that all models universally underperform when tested on Portuguese datasets. This is the direct result of the fact that data points for Portuguese almost exclusively represent the domain of social media. As a consequence, some combinations of filtering facets in our dataset collection produce very little data (i.e., asking for social media data in the Germanic genus of Indo-European languages will produce a significantly larger dataset than asking for news data representing Afro-Asiatic languages).

Finally, we acknowledge the lack of internal coherence of annotation protocols between datasets and languages. We have enforced strict quality criteria and rejected all datasets published without the annotation protocol, but we were unable, for obvious reasons, to unify annotation guidelines. The annotation of sentiment expressions and the assignment of sentiment labels are heavily subjective and, at the same time, influenced by cultural and linguistic features. Unfortunately, it is possible that semantically similar utterances will be assigned conflicting labels if they come from different datasets or modalities.

Acknowledgments and Disclosure of Funding

The work was partially supported by the (1) European Regional Development Fund (ERDF) in RPO WD 2014-2020 (project no. RPDS.01.02.01-02-0065/20); (2) the Polish Ministry of Education and Science, CLARIN-PL; (3) the ERDF in the 2014-2020 Smart Growth Operational Programme, CLARIN – Common Language Resources and Technology Infrastructure; (4) project CLARIN-Q (agreement no. 2022/WK/09); (5) the Department of Artificial Intelligence at Wroclaw University of Science and Technology.

References

[1] Marwan Al Omari, Moustafa Al-Hajj, Nacereddine Hammami, and Amani Sabra. Sentiment classifier: Logistic regression for arabic services' reviews in lebanon. In *2019 International Conference on Computer and Information Sciences (ICCIS)*, pages 1–5, 2019. doi: 10.1109/ICCISci.2019.8716394.

[2] Mohamed Aly and Amir Atiya. LABR: A large scale Arabic book reviews dataset. In *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 494–498, Sofia, Bulgaria, August 2013. Association for Computational Linguistics. URL `https://aclanthology.org/P13-2088`.

[3] Adam Amram, Anat Ben David, and Reut Tsarfaty. Representations and architectures in neural sentiment analysis for morphologically rich languages: A case study from Modern Hebrew. In *Proceedings of the 27th International Conference on Computational Linguistics*, pages 2242–2252, Santa Fe, New Mexico, USA, August 2018. Association for Computational Linguistics. URL `https://aclanthology.org/C18-1190`.

[4] Monika Arora and Vineet Kansal. Character level embedding with deep convolutional neural network for text normalization of unstructured data for twitter sentiment analysis. *Social Network Analysis and Mining*, 9(1):12, Mar 2019. ISSN 1869-5469. doi: 10.1007/s13278-019-0557-y. URL `https://doi.org/10.1007/s13278-019-0557-y`.

[5] Mikel Artetxe and Holger Schwenk. Massively Multilingual Sentence Embeddings for Zero-Shot Cross-Lingual Transfer and Beyond. *Transactions of the Association for Computational Linguistics*, 7:597–610, 09 2019. ISSN 2307-387X. doi: 10.1162/tacl_a_00288. URL `https://doi.org/10.1162/tacl_a_00288`.

[6] Mohammed Attia, Younes Samih, Ali Elkahky, and Laura Kallmeyer. Multilingual multi-class sentiment classification using convolutional neural networks. In *Proceedings of the Eleventh International Conference on Language Resources and Evaluation (LREC 2018)*, Miyazaki, Japan, May 2018. European Language Resources Association (ELRA). URL https://aclanthology.org/L18-1101.

[7] Ramy Baly, Alaa Khaddaj, Hazem M. Hajj, Wassim El-Hajj, and Khaled Bashir Shaban. ArSentD-LEV: A Multi-Topic Corpus for Target-based Sentiment Analysis in Arabic Levantine Tweets. In Hend Al-Khalifa, King Saud University, KSA Walid Magdy, University of Edinburgh, UK Kareem Darwish, Qatar Computing Research Institute, Qatar Tamer Elsayed, Qatar University, and Qatar, editors, *Proceedings of the Eleventh International Conference on Language Resources and Evaluation (LREC 2018)*, Paris, France, may 2018. European Language Resources Association (ELRA). ISBN 979-10-95546-25-2.

[8] Francesco Barbieri, Valerio Basile, Danilo Croce, Malvina Nissim, Nicole Novielli, and Viviana Patti. Overview of the Evalita 2016 SENTIment POLarity Classification Task. In *Proceedings of Third Italian Conference on Computational Linguistics (CLiC-it 2016) & Fifth Evaluation Campaign of Natural Language Processing and Speech Tools for Italian. Final Workshop (EVALITA 2016)*, Naples, Italy, December 2016. URL https://hal.inria.fr/hal-01414731.

[9] Francesco Barbieri, Luis Espinosa Anke, and Jose Camacho-Collados. Xlm-t: A multilingual language model toolkit for twitter. *arXiv e-prints*, pages arXiv–2104, 2021.

[10] Mohaddeseh Bastan, Mahnaz Koupaee, Youngseo Son, Richard Sicoli, and Niranjan Balasubramanian. Author's sentiment prediction. In *Proceedings of the 28th International Conference on Computational Linguistics*, pages 604–615, Barcelona, Spain (Online), December 2020. International Committee on Computational Linguistics. doi: 10.18653/v1/2020.coling-main.52. URL https://aclanthology.org/2020.coling-main.52.

[11] Vuk Batanović, Boško Nikolić, and Milan Milosavljević. Reliable baselines for sentiment analysis in resource-limited languages: The Serbian movie review dataset. In *Proceedings of the Tenth International Conference on Language Resources and Evaluation (LREC'16)*, pages 2688–2696, Portorož, Slovenia, May 2016. European Language Resources Association (ELRA). URL https://aclanthology.org/L16-1427.

[12] Vuk Batanović, Miloš Cvetanović, and Boško Nikolić. A versatile framework for resource-limited sentiment articulation, annotation, and analysis of short texts. *PLOS ONE*, 15(11):1–30, 11 2020. doi: 10.1371/journal.pone.0242050. URL https://doi.org/10.1371/journal.pone.0242050.

[13] Henrico Brum and Maria das Graças Volpe Nunes. Building a sentiment corpus of tweets in Brazilian Portuguese. In *Proceedings of the Eleventh International Conference on Language Resources and Evaluation (LREC 2018)*, Miyazaki, Japan, May 2018. European Language Resources Association (ELRA). URL https://aclanthology.org/L18-1658.

[14] Jože Bučar, Martin Žnidaršič, and Janez Povh. Annotated news corpora and a lexicon for sentiment analysis in slovene. *Language Resources and Evaluation*, 52(3):895–919, September 2018. doi: 10.1007/s10579-018-9413-3.

[15] Ethem F. Can, Aysu Ezen-Can, and Fazli Can. Multilingual sentiment analysis: An RNN-based framework for limited data. *Computing Research Repository*, arXiv:1806.04511, 2018. URL http://arxiv.org/abs/1806.04511. Version 1.

[16] Emile Chapuis, Pierre Colombo, Matteo Manica, Matthieu Labeau, and Chloé Clavel. Hierarchical pre-training for sequence labelling in spoken dialog. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 2636–2648, Online, November 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.findings-emnlp.239. URL https://aclanthology.org/2020.findings-emnlp.239.

[17] Mark Cieliebak, Jan Milan Deriu, Dominic Egger, and Fatih Uzdilli. A Twitter corpus and benchmark resources for German sentiment analysis. In *Proceedings of the Fifth International Workshop on Natural Language Processing for Social Media*, pages 45–51, Valencia, Spain,

April 2017. Association for Computational Linguistics. doi: 10.18653/v1/W17-1106. URL https://aclanthology.org/W17-1106.

[18] Alexis Conneau and Guillaume Lample. Cross-lingual language model pretraining. In *Proceedings of the 33rd International Conference on Neural Information Processing Systems*, Red Hook, NY, USA, 2019. Curran Associates Inc. URL https://dl.acm.org/doi/10.5555/3454287.3454921.

[19] Alexis Conneau, Ruty Rinott, Guillaume Lample, Adina Williams, Samuel Bowman, Holger Schwenk, and Veselin Stoyanov. XNLI: Evaluating cross-lingual sentence representations. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 2475–2485, Brussels, Belgium, October-November 2018. Association for Computational Linguistics. doi: 10.18653/v1/D18-1269. URL https://aclanthology.org/D18-1269.

[20] Alexis Conneau, Kartikay Khandelwal, Naman Goyal, Vishrav Chaudhary, Guillaume Wenzek, Francisco Guzmán, Edouard Grave, Myle Ott, Luke Zettlemoyer, and Veselin Stoyanov. Unsupervised cross-lingual representation learning at scale. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 8440–8451, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.747. URL https://aclanthology.org/2020.acl-main.747.

[21] Fermin L Cruz, Jose A Troyano, Fernando Enriquez, and Javier Ortega. Experiments in sentiment classification of movie reviews in spanish. *Procesamiento del Lenguaje Natural*, 41: 73–80, 2008.

[22] Kia Dashtipour, Soujanya Poria, Amir Hussain, Erik Cambria, Ahmad YA Hawalah, Alexander Gelbukh, and Qiang Zhou. Multilingual sentiment analysis: state of the art and independent comparison of techniques. *Cognitive computation*, 8(4):757–771, 2016. doi: 10.1007/s12559-016-9415-7.

[23] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1423. URL https://aclanthology.org/N19-1423.

[24] Matthew S. Dryer and Martin Haspelmath, editors. *WALS Online (v2020.3)*. Zenodo, 2013. doi: 10.5281/zenodo.7385533. URL https://doi.org/10.5281/zenodo.7385533.

[25] Ashraf Elnagar and Omar Einea. BRAD 1.0: Book reviews in arabic dataset. In *2016 IEEE/ACS 13th International Conference of Computer Systems and Applications (AICCSA)*, pages 1–8, 2016. doi: 10.1109/AICCSA.2016.7945800.

[26] Ashraf Elnagar, Yasmin S. Khalifa, and Anas Einea. *Hotel Arabic-Reviews Dataset Construction for Sentiment Analysis Applications*. Springer International Publishing, Cham, 2018. ISBN 978-3-319-67056-0. doi: 10.1007/978-3-319-67056-0_3. URL https://doi.org/10.1007/978-3-319-67056-0_3.

[27] Fangxiaoyu Feng, Yinfei Yang, Daniel Cer, Naveen Arivazhagan, and Wei Wang. Language-agnostic BERT Sentence Embedding. *Computing Research Repository*, arXiv:2007.01852, 2020. Version 2.

[28] Steven L Gordon. The sociology of sentiments and emotion. In *Social psychology*, pages 562–592. Routledge, 2017.

[29] Ivan Habernal, Tomáš Ptáček, and Josef Steinberger. Sentiment analysis in Czech social media using supervised machine learning. In *Proceedings of the 4th Workshop on Computational Approaches to Subjectivity, Sentiment and Social Media Analysis*, pages 65–74, Atlanta, Georgia, June 2013. Association for Computational Linguistics. URL https://aclanthology.org/W13-1609.

[30] Pedram Hosseini, Ali Ahmadian Ramaki, Hassan Maleki, Mansoureh Anvari, and Seyed Abolghasem Mirroshandel. SentiPers: A sentiment analysis corpus for persian. *Computing Research Repository*, arXiv:1801.07737, 2018. URL http://arxiv.org/abs/1801.07737. Version 2.

[31] Junjie Hu, Sebastian Ruder, Aditya Siddhant, Graham Neubig, Orhan Firat, and Melvin Johnson. XTREME: A massively multilingual multi-task benchmark for evaluating cross-lingual generalisation. In Hal Daumé III and Aarti Singh, editors, *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 4411–4421. PMLR, 13–18 Jul 2020. URL https://proceedings.mlr.press/v119/hu20b.html.

[32] Clayton J. Hutto and Eric Gilbert. VADER: A parsimonious rule-based model for sentiment analysis of social media text. In *Proceedings of the International AAAI Conference on Web and Social Media*, volume 8, pages 216–225, May 2014. URL https://ojs.aaai.org/index.php/ICWSM/article/view/14550.

[33] Crowdflower Inc. Twitter us airline sentiment, 2015. URL https://www.kaggle.com/crowdflower/twitter-airline-sentiment.

[34] Karthikeyan K, Zihan Wang, Stephen Mayhew, and Dan Roth. Cross-lingual ability of multilingual bert: An empirical study. In *International Conference on Learning Representations*, 2020. URL https://openreview.net/forum?id=HJeT3yrtDr.

[35] Kamil Kanclerz, Piotr Miłkowski, and Jan Kocoń. Cross-lingual deep neural transfer learning in sentiment analysis. *Procedia Computer Science*, 176:128–137, 2020. ISSN 1877-0509. doi: https://doi.org/10.1016/j.procs.2020.08.014. URL https://www.sciencedirect.com/science/article/pii/S187705092031838X. Knowledge-Based and Intelligent Information & Engineering Systems: Proceedings of the 24th International Conference KES2020.

[36] Brian Keith Norambuena, Exequiel Lettura, and Claudio Villegas. Sentiment analysis and opinion mining applied to scientific paper reviews. *Intelligent Data Analysis*, 23:191–214, 02 2019. doi: 10.3233/IDA-173807.

[37] Phillip Keung, Yichao Lu, György Szarvas, and Noah A. Smith. The multilingual Amazon reviews corpus. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 4563–4568, Online, November 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.emnlp-main.369. URL https://aclanthology.org/2020.emnlp-main.369.

[38] Jan Kocoń, Piotr Miłkowski, and Monika Zaśko-Zielińska. Multi-level sentiment analysis of PolEmo 2.0: Extended corpus of multi-domain consumer reviews. In *Proceedings of the 23rd Conference on Computational Natural Language Learning (CoNLL)*, pages 980–991, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/K19-1092. URL https://aclanthology.org/K19-1092.

[39] Patrick Lewis, Barlas Oguz, Ruty Rinott, Sebastian Riedel, and Holger Schwenk. MLQA: Evaluating cross-lingual extractive question answering. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 7315–7330, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.653. URL https://aclanthology.org/2020.acl-main.653.

[40] Jindřich Libovický, Rudolf Rosa, and Alexander Fraser. How language-neutral is multilingual bert? *Computing Research Repository*, arXiv:1911.03310, 2019. Version 1.

[41] Yiou Lin, Hang Lei, Jia Wu, and Xiaoyu Li. An empirical study on sentiment classification of Chinese review using word embedding. In *Proceedings of the 29th Pacific Asia Conference on Language, Information and Computation: Posters*, pages 258–266, Shanghai, China, October 2015. URL https://aclanthology.org/Y15-2030.

[42] Chi-Liang Liu, Tsung-Yuan Hsu, Yung-Sung Chuang, and Hung yi Lee. What makes multilingual bert multilingual? *Computing Research Repository*, arXiv:2010.10938, 2020. Version 1.

[43] Alexandre Magueresse, Vincent Carles, and Evan Heetderks. Low-resource languages: A review of past work and future challenges. *arXiv preprint arXiv:2006.07264*, 2020.

[44] Pekka Malo, Ankur Sinha, Pekka Korhonen, Jyrki Wallenius, and Pyry Takala. Good debt or bad debt: Detecting semantic orientations in economic texts. *Journal of the Association for Information Science and Technology*, 65(4):782–796, apr 2014. ISSN 2330-1635. doi: 10.1002/asi.23062. URL https://doi.org/10.1002/asi.23062.

[45] Tomás Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient estimation of word representations in vector space. In Yoshua Bengio and Yann LeCun, editors, *1st International Conference on Learning Representations, ICLR 2013, Scottsdale, Arizona, USA, May 2-4, 2013, Workshop Track Proceedings*, 2013. URL http://arxiv.org/abs/1301.3781.

[46] Tomas Mikolov, Quoc V. Le, and Ilya Sutskever. Exploiting similarities among languages for machine translation. *Computing Research Repository*, arXiv:1309.4168, 2013. Version 1.

[47] Igor Mozetič, Miha Grčar, and Jasmina Smailović. Multilingual twitter sentiment classification: The role of human annotators. *PLOS ONE*, 11(5):1–26, 05 2016. doi: 10.1371/journal.pone.0155036. URL https://doi.org/10.1371/journal.pone.0155036.

[48] Mahmoud Nabil, Mohamed Aly, and Amir Atiya. ASTD: Arabic sentiment tweets dataset. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*, pages 2515–2519, Lisbon, Portugal, September 2015. Association for Computational Linguistics. doi: 10.18653/v1/D15-1299. URL https://aclanthology.org/D15-1299.

[49] Sascha Narr, Michael Hülfenhaus, and Sahin Albayrak. Language-independent twitter sentiment analysis. In *Workshop on Knowledge Discovery, Data Mining and Machine Learning (KDML-2012)*, 2012.

[50] Jianmo Ni, Jiacheng Li, and Julian McAuley. Justifying recommendations using distantly-labeled reviews and fine-grained aspects. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 188–197, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-1018. URL https://aclanthology.org/D19-1018.

[51] Curtis Northcutt, Lu Jiang, and Isaac Chuang. Confident learning: Estimating uncertainty in dataset labels. *Journal of Artificial Intelligence Research*, 70:1373–1411, 2021.

[52] Parth Patwa, Gustavo Aguilar, Sudipta Kar, Suraj Pandey, Srinivas PYKL, Björn Gambäck, Tanmoy Chakraborty, Thamar Solorio, and Amitava Das. SemEval-2020 task 9: Overview of sentiment analysis of code-mixed tweets. In *Proceedings of the Fourteenth Workshop on Semantic Evaluation*, pages 774–790, Barcelona (online), December 2020. International Committee for Computational Linguistics. doi: 10.18653/v1/2020.semeval-1.100. URL https://aclanthology.org/2020.semeval-1.100.

[53] Parth Patwa, Gustavo Aguilar, Sudipta Kar, Suraj Pandey, Srinivas PYKL, Björn Gambäck, Tanmoy Chakraborty, Thamar Solorio, and Amitava Das. SemEval-2020 task 9: Overview of sentiment analysis of code-mixed tweets. In *Proceedings of the Fourteenth Workshop on Semantic Evaluation*, pages 774–790, Barcelona (online), December 2020. International Committee for Computational Linguistics. URL https://aclanthology.org/2020.semeval-1.100.

[54] Andraž Pelicon, Marko Pranjić, Dragana Miljković, Blaž Škrlj, and Senja Pollak. Zero-shot learning for cross-lingual news sentiment classification. *Applied Sciences*, 10(17):5993, 2020.

[55] Andraž Pelicon, Marko Pranjić, Dragana Miljković, Blaž Škrlj, and Senja Pollak. Zero-shot learning for cross-lingual news sentiment classification. *Applied Sciences*, 10(17), 2020. ISSN 2076-3417. doi: 10.3390/app10175993. URL https://www.mdpi.com/2076-3417/10/17/5993.

[56] Nils Reimers and Iryna Gurevych. Making monolingual sentence embeddings multilingual using knowledge distillation. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 4512–4525, Online, November 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.emnlp-main.365. URL https://aclanthology.org/2020.emnlp-main.365.

[57] Anna Rogers, Alexey Romanov, Anna Rumshisky, Svitlana Volkova, Mikhail Gronas, and Alex Gribov. RuSentiment: An enriched sentiment analysis dataset for social media in Russian. In *Proceedings of the 27th International Conference on Computational Linguistics*, pages 755–763, Santa Fe, New Mexico, USA, August 2018. Association for Computational Linguistics. URL https://aclanthology.org/C18-1064.

[58] Sara Rosenthal, Noura Farra, and Preslav Nakov. SemEval-2017 task 4: Sentiment analysis in Twitter. In *Proceedings of the 11th International Workshop on Semantic Evaluation (SemEval-2017)*, pages 502–518, Vancouver, Canada, August 2017. Association for Computational Linguistics. doi: 10.18653/v1/S17-2088. URL https://aclanthology.org/S17-2088.

[59] Sebastian Ruder, Ivan Vulić, and Anders Søgaard. A survey of cross-lingual word embedding models. *Journal of Artificial Intelligence Research*, 65:569–631, Aug 2019. ISSN 1076-9757. doi: 10.1613/jair.1.11640. URL http://dx.doi.org/10.1613/jair.1.11640.

[60] Piotr Rybak, Robert Mroczkowski, Janusz Tracz, and Ireneusz Gawlik. KLEJ: Comprehensive benchmark for Polish language understanding. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 1191–1201, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.111. URL https://aclanthology.org/2020.acl-main.111.

[61] Santwana Sagnika, Anshuman Pattanaik, Bhabani Shankar Prasad Mishra, and Saroj K Meher. A review on multi-lingual sentiment analysis by machine learning methods. *Journal of Engineering Science & Technology Review*, 13(2):154–166, 2020. doi: 10.25103/jestr.132.19.

[62] Mohammad Salameh, Saif Mohammad, and Svetlana Kiritchenko. Sentiment after translation: A case-study on Arabic social media posts. In *Proceedings of the 2015 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 767–777, Denver, Colorado, May–June 2015. Association for Computational Linguistics. doi: 10.3115/v1/N15-1078. URL https://aclanthology.org/N15-1078.

[63] Niek J Sanders. Sanders-Twitter Sentiment Corpus. *Sanders Analytics LLC*, 2011.

[64] Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter. *Computing Research Repository*, arXiv:1910.01108, 2020. Version 4.

[65] Anindya Sarkar, Sujeeth Reddy, and Raghu Sesha Iyengar. Zero-shot multilingual sentiment analysis using hierarchical attentive network and bert. In *Proceedings of the 2019 3rd International Conference on Natural Language Processing and Information Retrieval*, pages 49–56, 2019.

[66] Dietmar Schabus and Marcin Skowron. Academic-industrial perspective on the development and deployment of a moderation system for a newspaper website. In *Proceedings of the Eleventh International Conference on Language Resources and Evaluation (LREC 2018)*, Miyazaki, Japan, May 2018. European Language Resources Association (ELRA). URL https://aclanthology.org/L18-1253.

[67] Zareen Sharf and Saif Ur Rahman. Performing natural language processing on roman urdu datasets. In *International Journal of Computer Science and Network Security*, volume 18, pages 141–148, 2018. URL http://paper.ijcsns.org/07_book/201801/20180117.pdf.

[68] Emily Sheng and David Uthus. Investigating societal biases in a poetry composition system. In *Proceedings of the Second Workshop on Gender Bias in Natural Language Processing*, pages 93–106, Barcelona, Spain (Online), December 2020. Association for Computational Linguistics. URL https://aclanthology.org/2020.gebnlp-1.9.

[69] Jasdeep Singh, Bryan McCann, Richard Socher, and Caiming Xiong. BERT is not an interlingua and the bias of tokenization. In *Proceedings of the 2nd Workshop on Deep Learning Approaches for Low-Resource NLP (DeepLo 2019)*, pages 47–55, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-6106. URL https://aclanthology.org/D19-6106.

[70] Pawel Sobkowicz and Antoni Sobkowicz. Two-year study of emotion and communication patterns in a highly polarized political discussion forum. *Social Science Computer Review*, 30 (4):448–469, 2012. doi: 10.1177/0894439312436512.

[71] Uga Sprogis and Matiss Rikters. What can we learn from almost a decade of food tweets. *Computing Research Repository*, arXiv:2007.05194, 2020. URL https://arxiv.org/abs/2007.05194. Version 2.

[72] Rachele Sprugnoli. Multiemotions-it: a new dataset for opinion polarity and emotion analysis for italian. In *Proceedings of the Seventh Italian Conference on Computational Linguistics*, 12 2020.

[73] Arthit Suriyawongkul, Ekapol Chuangsuwanich, Pattarawat Chormai, and Charin Polpanumas. Pythainlp/wisesight-sentiment: First release (v1.0), September 2019. URL https://doi.org/10.5281/zenodo.3457447. Zenodo.

[74] Mike Thelwall, Kevan Buckley, and Georgios Paltoglou. Sentiment strength detection for the social web. *J. Am. Soc. Inf. Sci. Technol.*, 63(1):163–173, January 2012. ISSN 1532-2882. doi: 10.1002/asi.21662. URL https://doi.org/10.1002/asi.21662.

[75] Ekkalak Thongthanomkul, Tanapol Nearunchorn, and Yuwat Chuesathuchon. wongnai-corpus. https://github.com/wongnai/wongnai-corpus, 2019.

[76] Erik F. Tjong Kim Sang. Introduction to the CoNLL-2002 shared task: Language-independent named entity recognition. In *COLING-02: The 6th Conference on Natural Language Learning 2002 (CoNLL-2002)*, 2002. URL https://aclanthology.org/W02-2024.

[77] Erik F. Tjong Kim Sang and Fien De Meulder. Introduction to the CoNLL-2003 shared task: Language-independent named entity recognition. In *Proceedings of the Seventh Conference on Natural Language Learning at HLT-NAACL 2003*, pages 142–147, 2003. URL https://aclanthology.org/W03-0419.

[78] Niels van der Heijden, Samira Abnar, and Ekaterina Shutova. A comparison of architectures and pretraining methods for contextualized multilingual word embeddings. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 9090–9097, 2020.

[79] Joonatas Wehrmann, Willian Becker, Henry E. L. Cagnini, and Rodrigo C. Barros. A character-based convolutional neural network for language-agnostic twitter sentiment analysis. In *2017 International Joint Conference on Neural Networks (IJCNN)*, pages 2384–2391, 2017. doi: 10.1109/IJCNN.2017.7966145.

[80] Shijie Wu and Mark Dredze. Are all languages created equal in multilingual BERT? In *Proceedings of the 5th Workshop on Representation Learning for NLP*, pages 120–130, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.repl4nlp-1.16. URL https://aclanthology.org/2020.repl4nlp-1.16.

[81] Min Xiao and Yuhong Guo. Distributed word representation learning for cross-lingual dependency parsing. In *Proceedings of the Eighteenth Conference on Computational Natural Language Learning*, pages 119–129, Ann Arbor, Michigan, June 2014. Association for Computational Linguistics. doi: 10.3115/v1/W14-1613. URL https://aclanthology.org/W14-1613.

[82] Linting Xue, Noah Constant, Adam Roberts, Mihir Kale, Rami Al-Rfou, Aditya Siddhant, Aditya Barua, and Colin Raffel. mT5: A massively multilingual pre-trained text-to-text transformer. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 483–498, Online, June 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.naacl-main.41. URL https://aclanthology.org/2021.naacl-main.41.

[83] Yinfei Yang, Gustavo Hernandez Abrego, Steve Yuan, Mandy Guo, Qinlan Shen, Daniel Cer, Yun-hsuan Sung, Brian Strope, and Ray Kurzweil. Improving multilingual sentence embedding using bi-directional dual encoder with additive margin softmax. In *Proceedings of the Twenty-Eighth International Joint Conference on Artificial Intelligence, IJCAI-19*, pages 5370–5378. International Joint Conferences on Artificial Intelligence Organization, 7 2019. doi: 10.24963/ijcai.2019/746. URL https://doi.org/10.24963/ijcai.2019/746.

[84] Yinfei Yang, Daniel Cer, Amin Ahmad, Mandy Guo, Jax Law, Noah Constant, Gustavo Hernandez Abrego, Steve Yuan, Chris Tar, Yun-hsuan Sung, Brian Strope, and Ray Kurzweil. Multilingual universal sentence encoder for semantic retrieval. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics: System Demonstrations*, pages 87–94, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-demos.12. URL https://aclanthology.org/2020.acl-demos.12.

[85] Michał Ziemski, Marcin Junczys-Dowmunt, and Bruno Pouliquen. The united nations parallel corpus v1.0. In *Proceedings of the Tenth International Conference on Language Resources and Evaluation (LREC 2016)*, pages 3530–3534, Portorož, Slovenia, May 2016. European Language Resources Association (ELRA). URL https://www.aclweb.org/anthology/L16-1561.

[86] Pierre Zweigenbaum, Serge Sharoff, and Reinhard Rapp. Overview of the second BUCC shared task: Spotting parallel sentences in comparable corpora. In *Proceedings of the 10th Workshop on Building and Using Comparable Corpora*, pages 60–67, Vancouver, Canada, August 2017. Association for Computational Linguistics. doi: 10.18653/v1/W17-2512. URL https://www.aclweb.org/anthology/W17-2512.

A Datasets

We present detailed lists of datasets included in our research in Tables 3 and 4. They include language, category, dataset size, class balance, and basic dataset characteristics.

Table 3: List of all monolingual datasets used in experiments. Category (Cat.): R - Reviews, SM - Social Media, C - Chats, N - News, P - Poems, M - Mixed. HL - human labeled, #Words and #Chars are mean values

Paper	Lang	Cat.	HL	Samples	NEG/NEU/POS	#Words	#Char.
Al Omari et al. [1]	ar	R	No	3096	13.0/10.2/76.8	9	51
Elnagar et al. [26]	ar	R	No	400101	13.0/19.9/67.1	22	127
Aly and Atiya [2]	ar	R	No	6250	11.6/17.9/70.5	65	343
Elnagar and Einea [25]	ar	R	No	504007	15.4/21.0/63.6	77	424
Baly et al. [7]	ar	SM	Yes	2809	47.2/23.9/29.0	22	130
Nabil et al. [48]	ar	SM	Yes	3224	50.9/25.0/24.1	16	94
Salameh et al. [62]	ar	SM	Yes	1199	48.0/10.5/41.5	11	51
Salameh et al. [62]	ar	SM	Yes	1998	67.5/10.1/22.4	20	107
Habernal et al. [29]	cs	R	No	91140	32.4/33.7/33.9	50	311
Habernal et al. [29]	cs	R	No	92758	7.9/23.4/68.7	20	131
Habernal et al. [29]	cs	SM	Yes	9752	20.4/53.1/26.5	10	59
Habernal et al. [29]	cs	SM	Yes	2637	30.8/60.6/8.6	33	170
Cieliebak et al. [17]	de	SM	Yes	9948	16.3/59.2/24.6	11	86
Schabus and Skowron [66]	de	SM	Yes	3598	47.3/51.5/1.2	33	237
Chapuis et al. [16]	en	C	Yes	12138	31.8/46.5/21.7	12	48
Chapuis et al. [16]	en	C	Yes	4643	22.3/48.9/28.8	15	71
Malo et al. [44]	en	N	Yes	3448	12.2/62.1/25.7	22	124
Bastan et al. [10]	en	N	Yes	5333	11.6/37.3/51.0	388	2129
Hutto and Gilbert [32]	en	N	No	5190	29.3/52.9/17.8	17	104
Sheng and Uthus [68]	en	P	Yes	1052	18.3/15.8/65.9	7	37
Hutto and Gilbert [32]	en	R	No	3708	34.2/19.5/46.3	16	87
Hutto and Gilbert [32]	en	R	No	10605	49.6/1.5/48.9	19	111
Ni et al. [50]	en	R	No	1883238	8.3/8.0/83.7	70	382
Sanders [63]	en	SM	Yes	3424	16.7/68.1/15.2	14	97
Thelwall et al. [74]	en	SM	Yes	11759	28.0/34.0/38.0	26	147
Inc. [33]	en	SM	Yes	14427	63.0/21.2/15.8	17	104
Hutto and Gilbert [32]	en	SM	No	4200	26.9/17.0/56.1	13	79
Keith Norambuena et al. [36]	es	R	Yes	382	45.0/27.2/27.8	165	1033
Cruz et al. [21]	es	R	No	3871	32.9/32.3/34.9	511	3000
Hosseini et al. [30]	fa	R	Yes	13525	12.0/37.5/50.5	21	104
Amram et al. [3]	he	SM	Yes	8619	26.5/2.8/70.8	22	110
Pelicon et al. [55]	hr	N	Yes	2025	22.5/61.4/16.0	161	1021
Barbieri et al. [8]	it	SM	Yes	8926	36.7/41.7/21.6	14	101
Sprugnoli [72]	it	SM	Yes	3139	24.4/14.9/60.6	17	106
Sprogis and Rikters [71]	lv	SM	Yes	5790	23.8/45.2/31.0	20	138
Rybak et al. [60]	pl	R	No	10074	30.8/13.2/56.0	80	494
Kocoń et al. [38]	pl	R	Yes	57038	42.4/26.8/30.8	30	175
Sobkowicz and Sobkowicz [70]	pl	SM	Yes	645	50.7/47.3/2.0	33	230
Brum and Volpe Nunes [13]	pt	SM	Yes	10109	28.8/25.1/46.1	12	74
Rogers et al. [57]	ru	SM	Yes	23226	16.8/54.6/28.6	12	79
Bučar et al. [14]	sl	N	Yes	10417	32.0/52.0/16.0	309	2017
Batanović et al. [11]	sr	R	No	4724	17.8/43.7/38.5	498	3097
Batanović et al. [12]	sr	R	Yes	3948	30.3/18.1/51.5	18	105
Thongthanomkul et al. [75]	th	R	No	46193	5.4/30.5/64.1	29	544
Suriyawongkul et al. [73]	th	SM	Yes	26126	26.1/55.6/18.3	6	90
Sharf and Rahman [67]	ur	M	Yes	19660	26.7/43.6/29.7	13	69
Lin et al. [41]	zh	R	No	125725	28.6/21.9/49.5	51	128

Table 4: List of all multilingual datasets used in experiments. Category (Cat.): R - Reviews, SM - Social Media, C - Chats, N - News, P - Poems, M - Mixed. HL - human labeled

Paper	Cat.	Lang	HL	Samples	(NEG/NEU/POS)	#Words	#Char.
Narr et al. [49]	SM	de	Yes	953	10.0/75.1/14.9	12	80
		de	Yes	1781	16.9/63.3/19.8	13	81
		en	Yes	7073	17.4/60.0/22.6	14	78
		fr	Yes	685	23.4/53.4/23.2	14	82
		fr	Yes	1786	25.0/54.3/20.8	15	83
		pt	Yes	759	28.1/33.2/38.7	14	78
		pt	Yes	1769	30.7/33.9/35.4	14	78
Keung et al. [37]	R	de	No	209073	40.1/20.0/39.9	33	208
		en	No	209393	40.0/20.0/40.0	34	179
		es	No	208127	40.2/20.0/39.8	27	152
		fr	No	208160	40.2/20.1/39.7	28	160
		ja	No	209780	40.0/20.0/40.0	2	101
		zh	No	205977	39.8/20.1/40.1	1	50
Rosenthal et al. [58]	M	ar	Yes	9391	35.5/40.6/23.9	14	105
		en	Yes	65071	19.1/45.7/35.2	18	111
Patwa et al. [52]	SM	es	Yes	14920	16.8/33.1/50.0	16	86
		hi	Yes	16999	29.4/37.6/33.0	27	128
Mozetič et al. [47]	SM	bg	Yes	62150	22.6/45.9/31.5	12	85
		bs	Yes	36183	33.4/30.5/36.1	12	75
		de	Yes	90534	19.7/52.8/27.4	12	94
		en	Yes	85784	26.8/44.1/29.1	12	77
		es	Yes	191412	11.8/37.9/50.3	14	92
		hr	Yes	75569	25.7/23.9/50.4	12	91
		hu	Yes	56682	15.9/31.0/53.1	11	83
		pl	Yes	168931	30.0/26.1/43.9	11	82
		pt	Yes	145197	37.2/35.0/27.8	10	61
		ru	Yes	87704	32.0/40.1/27.8	10	67
		sk	Yes	56623	25.6/22.5/51.9	13	97
		sl	Yes	103126	29.9/43.3/26.8	13	91
		sq	Yes	44284	15.7/33.1/51.1	13	90
		sr	Yes	67696	34.8/42.8/22.4	13	81
		sv	Yes	41346	40.3/31.2/28.5	14	94

B Experiments

B.1 Data quality

In this section, we expand upon the results of experiments included in the benchmark. We begin with the presentation of the distribution of data point quality across all datasets. The data point quality has been computed using the `cleanlab` library self-confidence label-quality score [51]. Figure 3 presents the cumulative distribution of data quality score. As we can see, the data quality still is a pressing issue, despite our efforts to publish a well-curated collection. Over 40% of all data points have quality 0.6 or worse. As we have mentioned, low-quality data points are not distributed uniformly across all languages and datasets, but the impact of low-quality data is most pronounced for low-resource languages.

B.2 Pairwise ranked comparison of models

Of course, everyone is most interested in finding the answer to the following question: "Which model should I use?" The answer, as usual, is: "It depends". Figure 4 presents the results of the Nemenyi pairwise rank comparison test. In neither of the considered training scenarios (HL, HB, TF), a single model is statistically better than others. However, these results allow us to draw partial conclusions. For instance, the MPNet performs best in the HL scenario, and the same model is not significantly worse than XML-M, which is the best model in HB and TF scenarios. We can also notice that mBERT-based models (mBERT and DistillmBERT) proved to be the worst language models for our

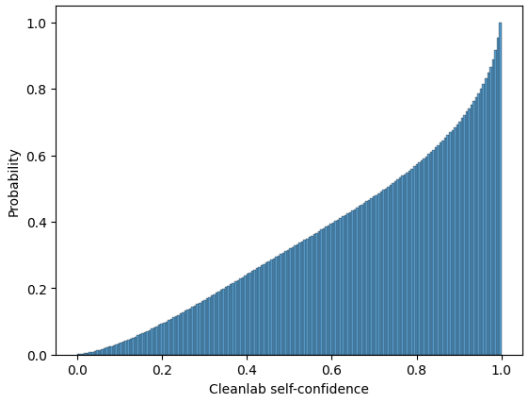

Figure 3: Cumulative distribution of sample quality across all datasets

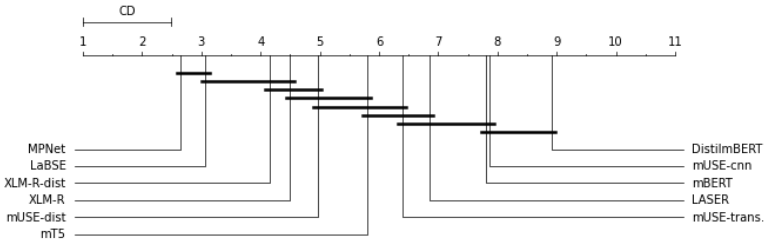

(a) HL scenario

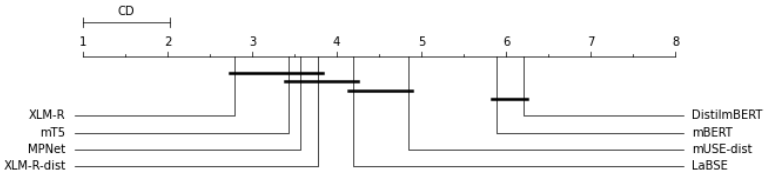

(b) HB scenario

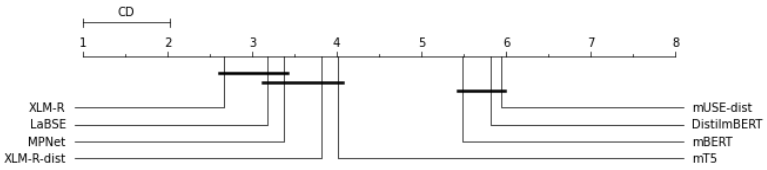

(c) FT scenario

Figure 4: Nemenyi diagrams of model ranks according to the F1-score on each dataset

tasks. Three models stand out as the most promising choices: XLM-R, LaBSE, and MPNet. They achieve comparable performance in all scenarios and test cases. Furthermore, they are better than other models in almost all test cases.

Table 5: Aggregated F_1 results of models. The best results for each test set are highlighted. W - whole test, A - avg. by dataset, I - internal

	XLM-R	LaBSE	MPNet	XLM-R-dist	mT5	mBERT	DistilmBERT	mUSE-dist	LASER	mUSE-trans.	mUSE-cnn
	HL scenario										
W	62	62	**63**	60	59	56	55	59	55	55	54
A	51	54	**55**	51	49	45	43	50	47	47	45
I	55	**61**	**61**	56	50	43	38	60	50	49	50
	HB scenario										
W	**66**	62	63	62	65	60	59	62	-	-	-
A	**57**	55	56	54	56	49	48	54	-	-	-
I	**64**	63	**64**	63	63	54	48	**64**	-	-	-
	FT scenario										
W	**68**	**68**	67	67	66	65	**64**	63	-	-	-
A	61	**62**	**62**	**62**	60	56	56	56	-	-	-
I	**70**	69	65	67	67	57	58	60	-	-	-

B.3 Detailed results for each model

Table 5 shows the F_1 scores of all models aggregated by datasets. The models improve with fine-tuning (up to 0.7 F_1) as compared to linear (0.61) or BiLSTM (0.64) layers operating on embeddings produced by models. The performance gains are more pronounced for models trained with masked language modeling objectives (mBERT, XML-R) than for models trained with sentence classification or sentence similarity objectives (LaBSE). Fine-tuning reduces inequalities in the results between models (0.55 vs. 0.43 for best and worst models in the HL scenario). The additional BiLSTM layer on top of transformer-based token embeddings has the capacity to improve the results compared to the linear classification layer model. The differences are most clear in the case of the results for our internal dataset, where the result improved even by 13 pp. for the mT5 model.

B.4 Detailed results for each language

We assessed the performance of each model in each experimental scenario per language. The texts were sub-sampled with stratification by language and class label so that language distribution in the test dataset follows the distribution in the whole dataset. We also include the total macro F_1 score value in column "all". Results are presented in Figure 5. Those results confirm our previous findings regarding the advantage of XLM-R, LaBSE, and MPNet models. They outperform other models in most languages, with no clear indication of superiority among them.

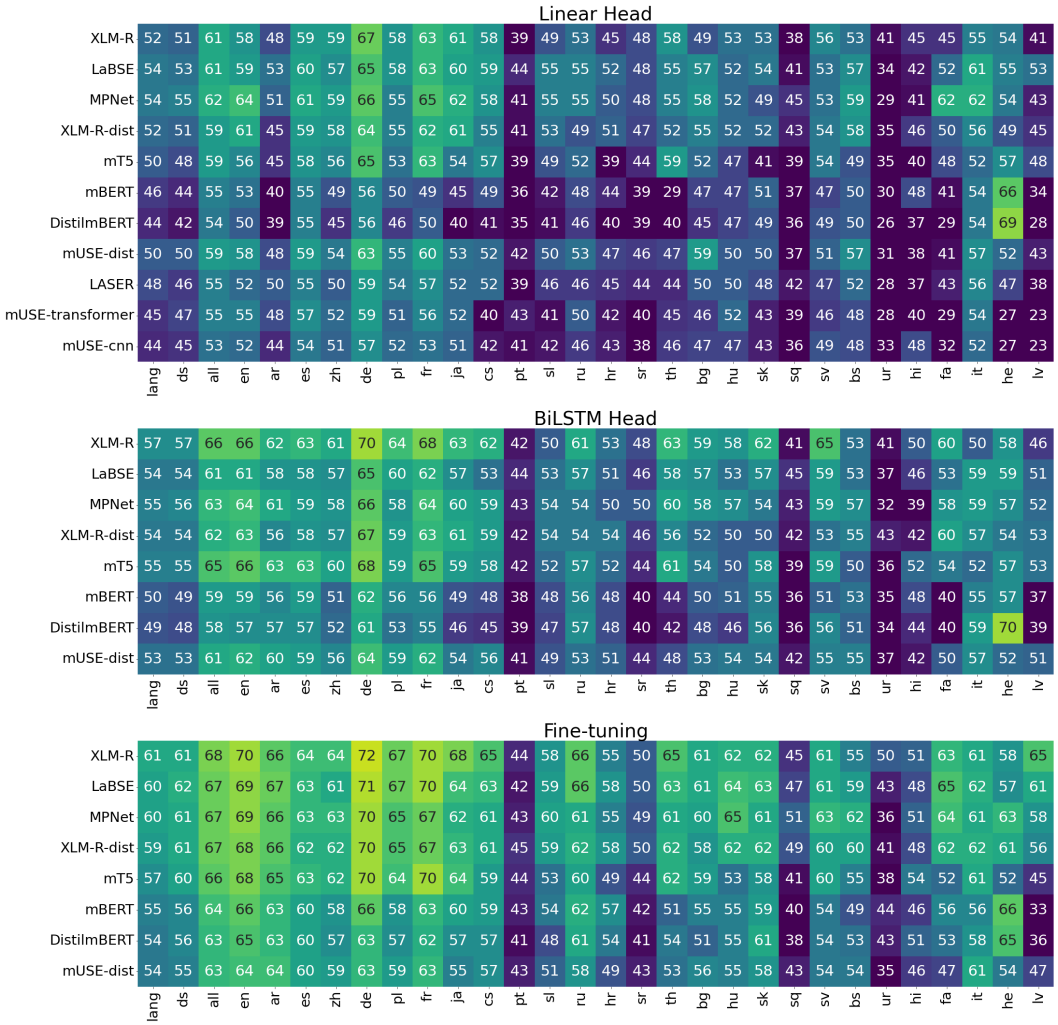

Figure 5: Detailed results of models' comparison. Legend: lang - averaged by all languages, ds - averaged by dataset, ar - Arabic, bg - Bulgarian, bs - Bosnian, cs - Czech, de - German, en - English, es - Spanish, fa - Persian, fr - French, he - Hebrew, hi - Hindi, hr - Croatian, hu - Hungarian, it - Italian, ja - Japanese, lv - Latvian, pl - Polish, pt - Portuguese, ru - Russian, sk - Slovak, sl - Slovenian, sq - Albanian, sr - Serbian, sv - Swedish, th - Thai, ur - Urdu, zh - Chinese.

B.5 Transfer between data modalities

The last example compares the effectiveness of transfer learning between data modalities. The results are presented in Figure 6. As expected, when models are tested against datasets from the same domain (news, social media, reviews), the average performance is much higher than on out-of-domain datasets. What is striking is the visible deterioration of the performance of models trained in the news domain. When knowledge transfer happens between social media and review domains, the average performance of models stays relatively the same.

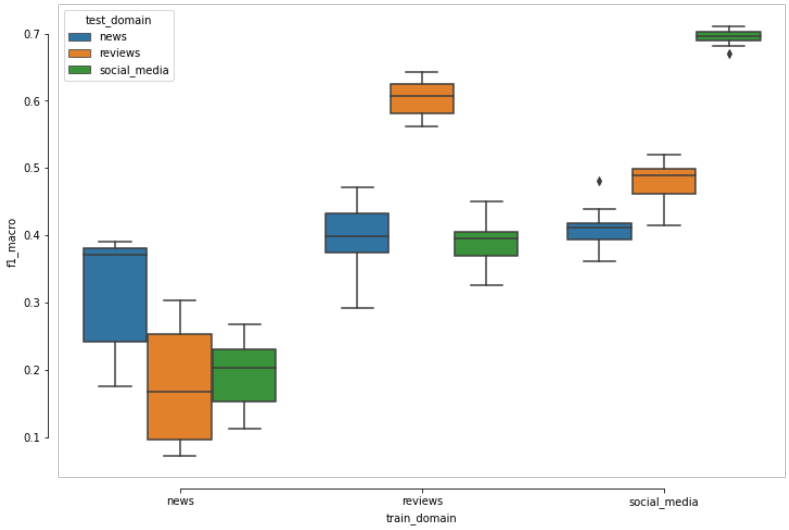

Figure 6: Transfer learning between data modalities

C Listings

The ease of loading and using more than 6M sentiment annotated texts is as simple as the listings above. All data will be downloaded automatically; you must only appropriately filter it to your use case.

```
import datasets

mms_dataset = datasets.load_dataset("Brand24/mms")
slavic = mms_dataset.filter(lambda row: row["Genus"] == "Slavic" and
    row["Polar questions"] == "interrogative word order")
```

Listing 1: Loading and filtering sentiment examples for Slavic *genus* with polar questions formed using interrogative word order.

```
import datasets

mms_dataset = datasets.load_dataset("Brand24/mms")
afro_asiatic = mms_dataset.filter(lambda row: row["Family"] == "Afro-
    Asiatic" and row["Number of cases"] == "no morphological case-
    making")
```

Listing 2: Loading and filtering sentiment examples for Afro-Asiatic Language Family with no morphological case-making.